

OPTIMIZATION OF A FLIP ALGORITHM

ETHZ – ADVANCED SYSTEM LAB PROJECT

CHRIS AMEVOR

SEAN BONE


FELIX ILLES

MIKAEL STELLIO

JUNE 7, 2021

THE FLIP ALGORITHM – INTRODUCTION

- FLIP: **F**luid **I**mplicit **P**article
- **Hybrid** particle- and grid-based method
- Commonly used in computer graphics to simulate fluids




THE FLIP ALGORITHM – COMPONENTS

At each timestep it

1. Computes the velocity field by **particle-to-grid** projection,
2. Applies external **forces** to the velocity field,
3. Enforces **boundary conditions**,
4. Computes **pressure** gradients and updates the velocity field,
5. Updates particle velocities using **grid-to-particle** projection,
6. **Advects** particles using 2nd-order Runge-Kutta integration,
7. Computes a **level set** & mesh around particles.

OPTIMIZATIONS

INITIAL RUNTIME




OPTIMIZATIONS – LEVEL SET COMPUTATION

The major benefits came from:

- Skipping construction of Eigen vectors: work directly on [arrays](#).
- Providing `Eigen::Map` objects where needed (Marching Cubes implementation).
- Computing boundary cells separately, simplifying inner loop.
- [Inlining](#) methods.
- Arithmetic optimizations & strength reduction: skip square root by working with $|x|^2$, constant div \rightarrow mul.

RUNTIME – LEVEL SET



OPTIMIZATIONS – PARTICLES DATA STRUCTURE

- Less bandwidth wasted
- Improved locality
- No "compiler black box" effect

Array of structs

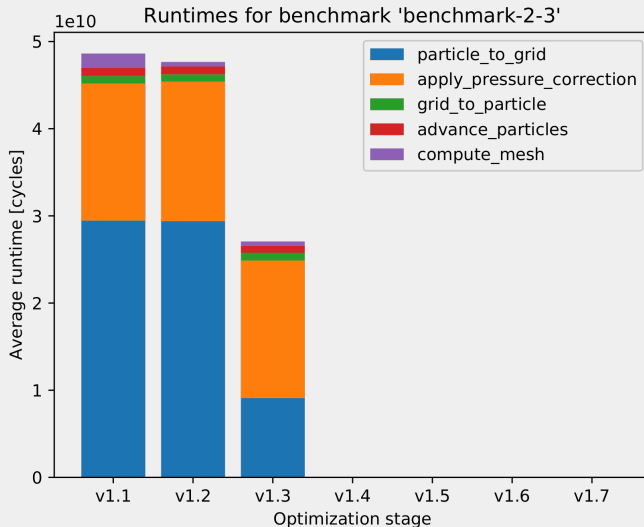


Struct of arrays

Two more optimizations were attempted without success (so far):

- **Particles sorting:** improves [locality](#), but causes overhead!
- **Cell index caching:** fewer flops at cost of more memory.

RUNTIME – PARTICLES DATA STRUCTURE



OPTIMIZATIONS – ADVECTION AND GRID-TO-PARTICLE

The main performance gain for both of these was in rewriting the [velocity interpolation routine](#):


- Fewer function calls.
- Simpler logic and less branching.
- [Template expressions](#) to "generate" different versions for U, V, and W.
- Fused U and U^* velocities interpolation in a single call. This again was done with templates and [benefits grid-to-particle](#).
- Strength reduction: dividing by cell size becomes multiplication.

OPTIMIZATIONS – ADVECTION AND GTP (SIMD)

Manual vectorization is **not beneficial**:

- Vectorizing just the trilinear interpolation kernel manually is a performance loss.
- Vectorizing the whole interpolation routine is infeasible due to complex branching (18 branches!).
- Vectorizing advection or grid-to-particle methods might be possible; however, they are dominated by interpolation.
- Because of the low footprint of advection and grid-to-particle, further optimization is deemed **unnecessary**.

RUNTIME – ADVECTION AND GRID-TO-PARTICLE



OPTIMIZATIONS – PARTICLE-TO-GRID

The major benefits came from:

- Manually **inlining** functions to allow loop fusion,
- Separating the accumulation loop for boundary cells,
- **Minimizing conditionals** inside loops.

Velocity fields U , V , and W have different dimensions!



Large overhead to avoid out-of-range accesses
(exploit ghost cells?)


OPTIMIZATIONS – PARTICLE-TO-GRID (SIMD)

Not very prone to vectorization and unrolling, due to:

- **Unordered** particles on the grid
- **Unknown number** of particles in a grid-cell

However, the **normalization** of accumulated velocities was completely vectorized resulting in a small speedup!

RUNTIME – PARTICLE-TO-GRID



OPTIMIZATIONS – PRESSURE CORRECTION

Original implementation using [sparse Eigen solver](#), requiring construction of sparse s.p.s.d. matrix A, where


$$A_{i,j} = A_{j,i} = \begin{cases} n_i & i = j, \\ -1 & \text{cells } i \text{ and } j \text{ adjacent,} \\ 0 & \text{otherwise.} \end{cases}$$

A custom **ICCG solver** allowed for better performance:

- No explicit representation of A , significantly **decreased memory usage** & bandwidth.
- Fused loops where possible.
- **Vectorization** of kernels.


OPTIMIZATIONS – PRESSURE CORRECTION

Blocking of forward & backward substitution to increase ILP was not beneficial.



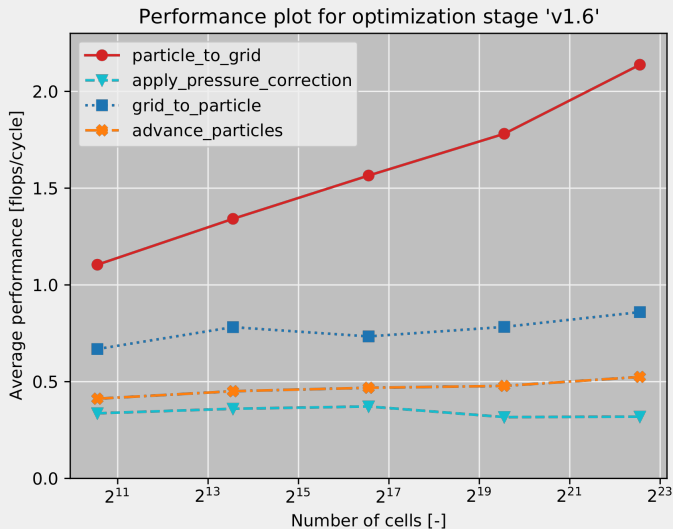
Problem: increases possible ILP at the cost of higher memory bandwidth!

RUNTIME – PRESSURE CORRECTION

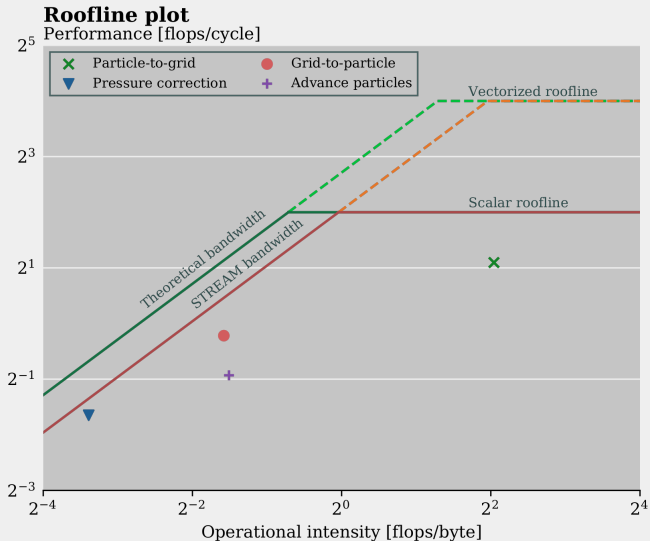


RESULTS

RESULTS – PERFORMANCE



RESULTS – ROOFLINE



THANKS FOR YOUR ATTENTION!